

*Lecture15-16.2 – OpenStack & Comparing
and Contrasting AWS with NeCTAR Cloud*

Professor Richard O. Sinnott & Yao Pan

University of Melbourne

rsinnott@unimelb.edu.au



&



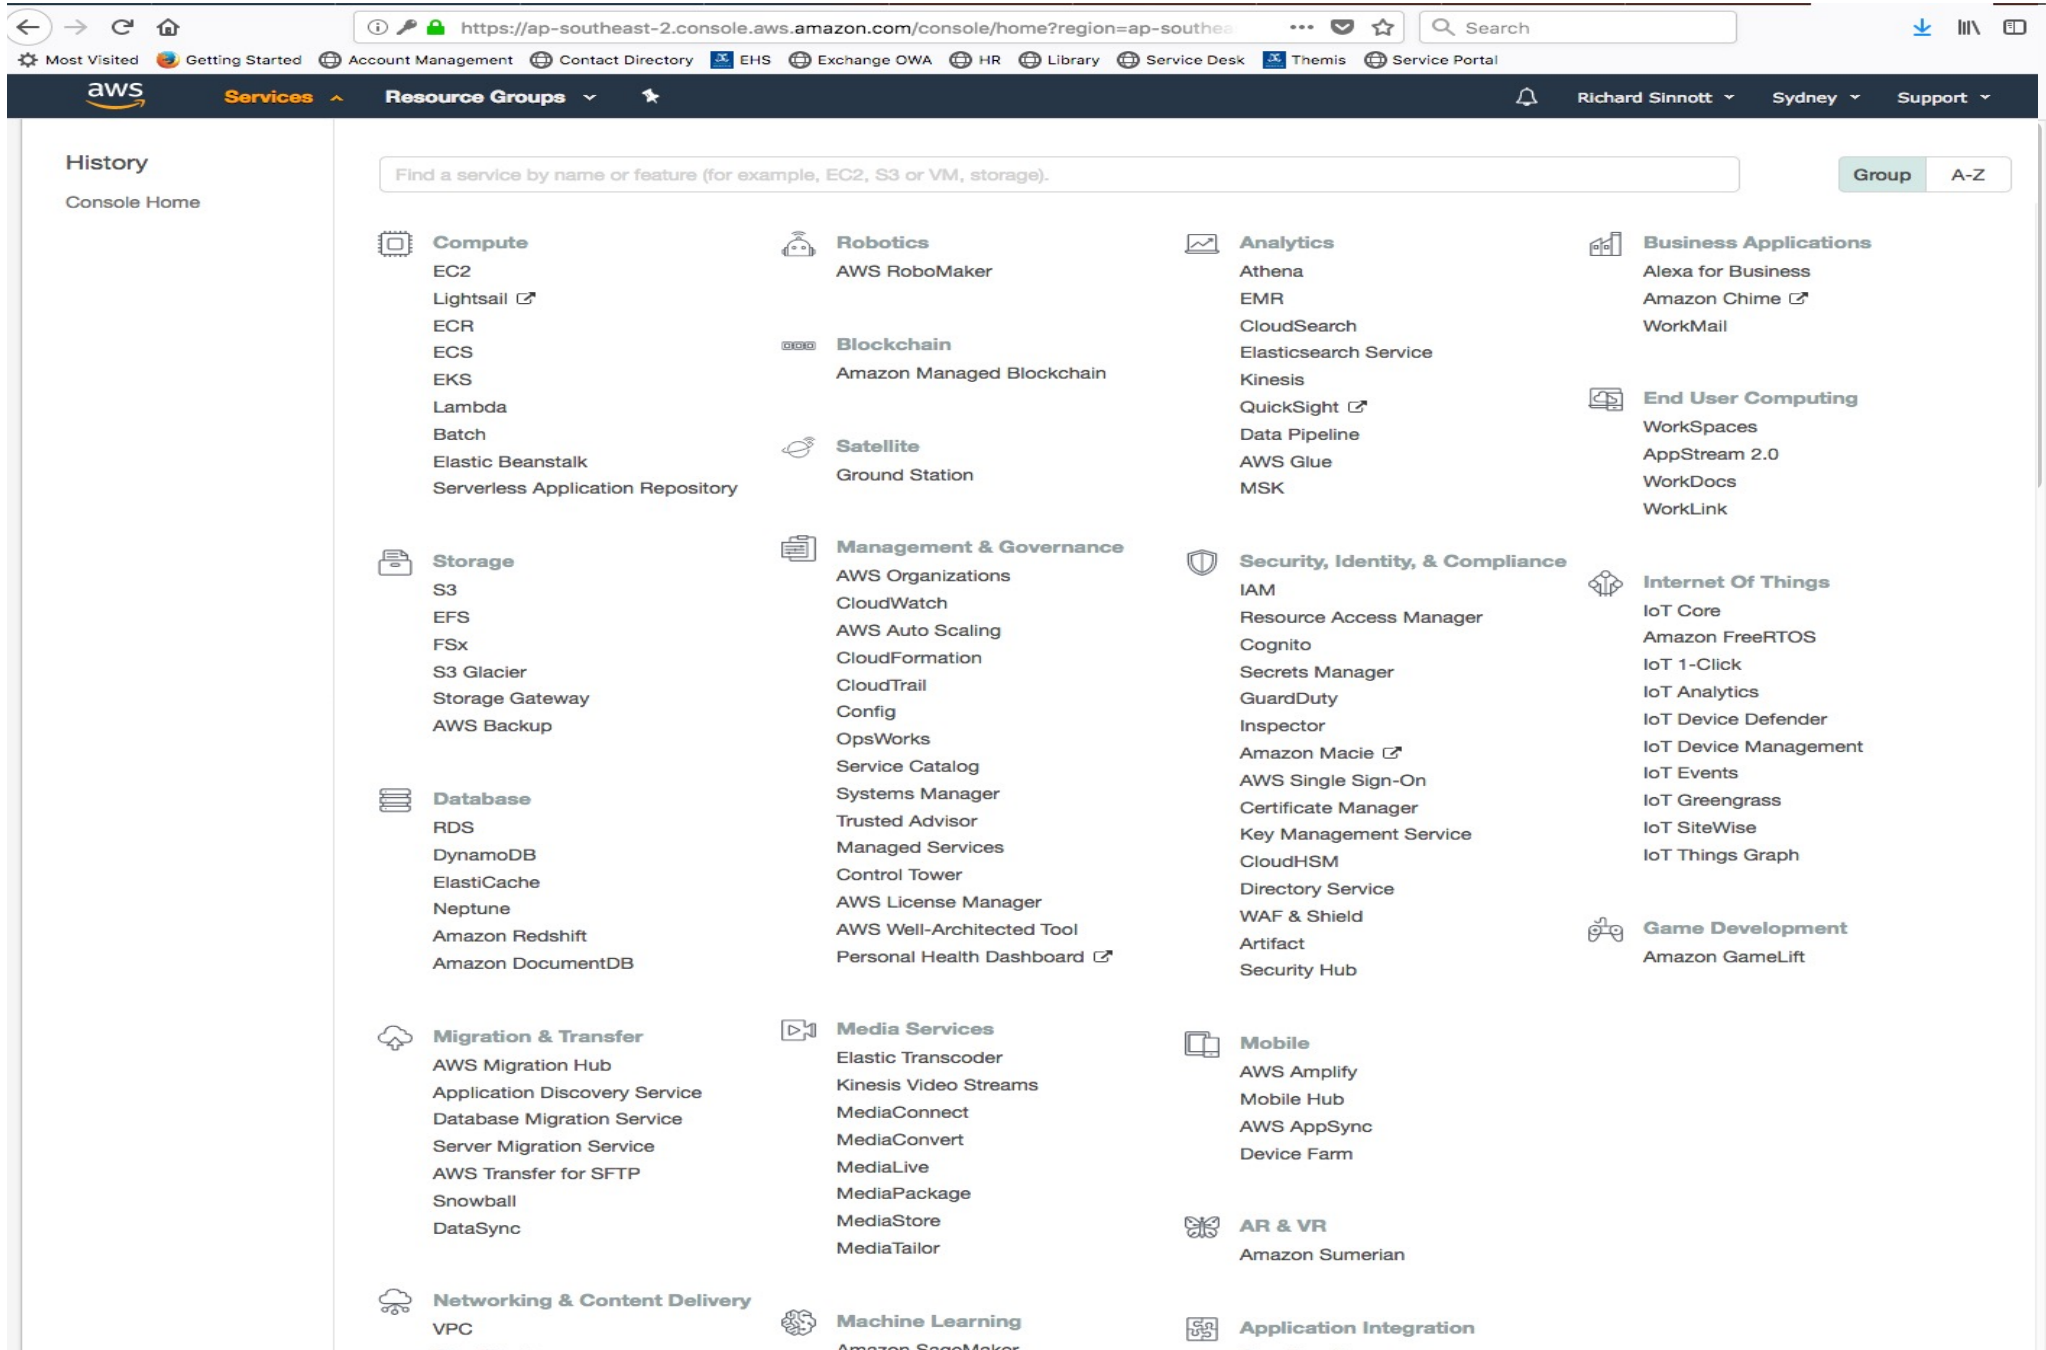
- UniMelb/NeCTAR Research Cloud
 - open source Cloud platform
 - openStack
 - overview of the major services*
- AWS (<http://aws.amazon.com>)
 - mainstream Cloud platform
 - Examples of the kinds of services that are available

*note that not all openstack services are available (yet!?) on the MRC/NeCTAR Research Cloud



nectar

&





nectar

OpenStack



- Began in 2010 as a joint project between Rackspace and NASA
- Offers free and open-source software platform for cloud computing for (mostly) IaaS
- Consists of interrelated components (services) that control / support compute, storage, and networking resources
- Often used through web-based dashboards, through command-line tools, or programmatically through ReSTful APIs
- Released under the terms of the Apache License
- Managed/coordinated by the OpenStack Foundation
 - non-profit corporate entity established in 2012 to promote OpenStack software and its community
 - Over 500 companies have since joined the project

OpenStack Components

- Many associated/underpinning services
 - Compute Service (code-named **Nova**)
 - Image Service (code-named **Glance**)
 - Block Storage Service (code named **Cinder**)
 - Object Storage Service (code-named **Swift**)
 - Security Management (code-named **Keystone**)
 - Orchestration Service (code-named **Heat**)
 - Network Service (code-named **Neutron**)
 - Container Service (code-named **Zun**)
 - Database service (code-named **Trove**)
 - Dashboard service (code-named **Horizon**)
 - Search service (code-named **Searchlight**)
 - ...

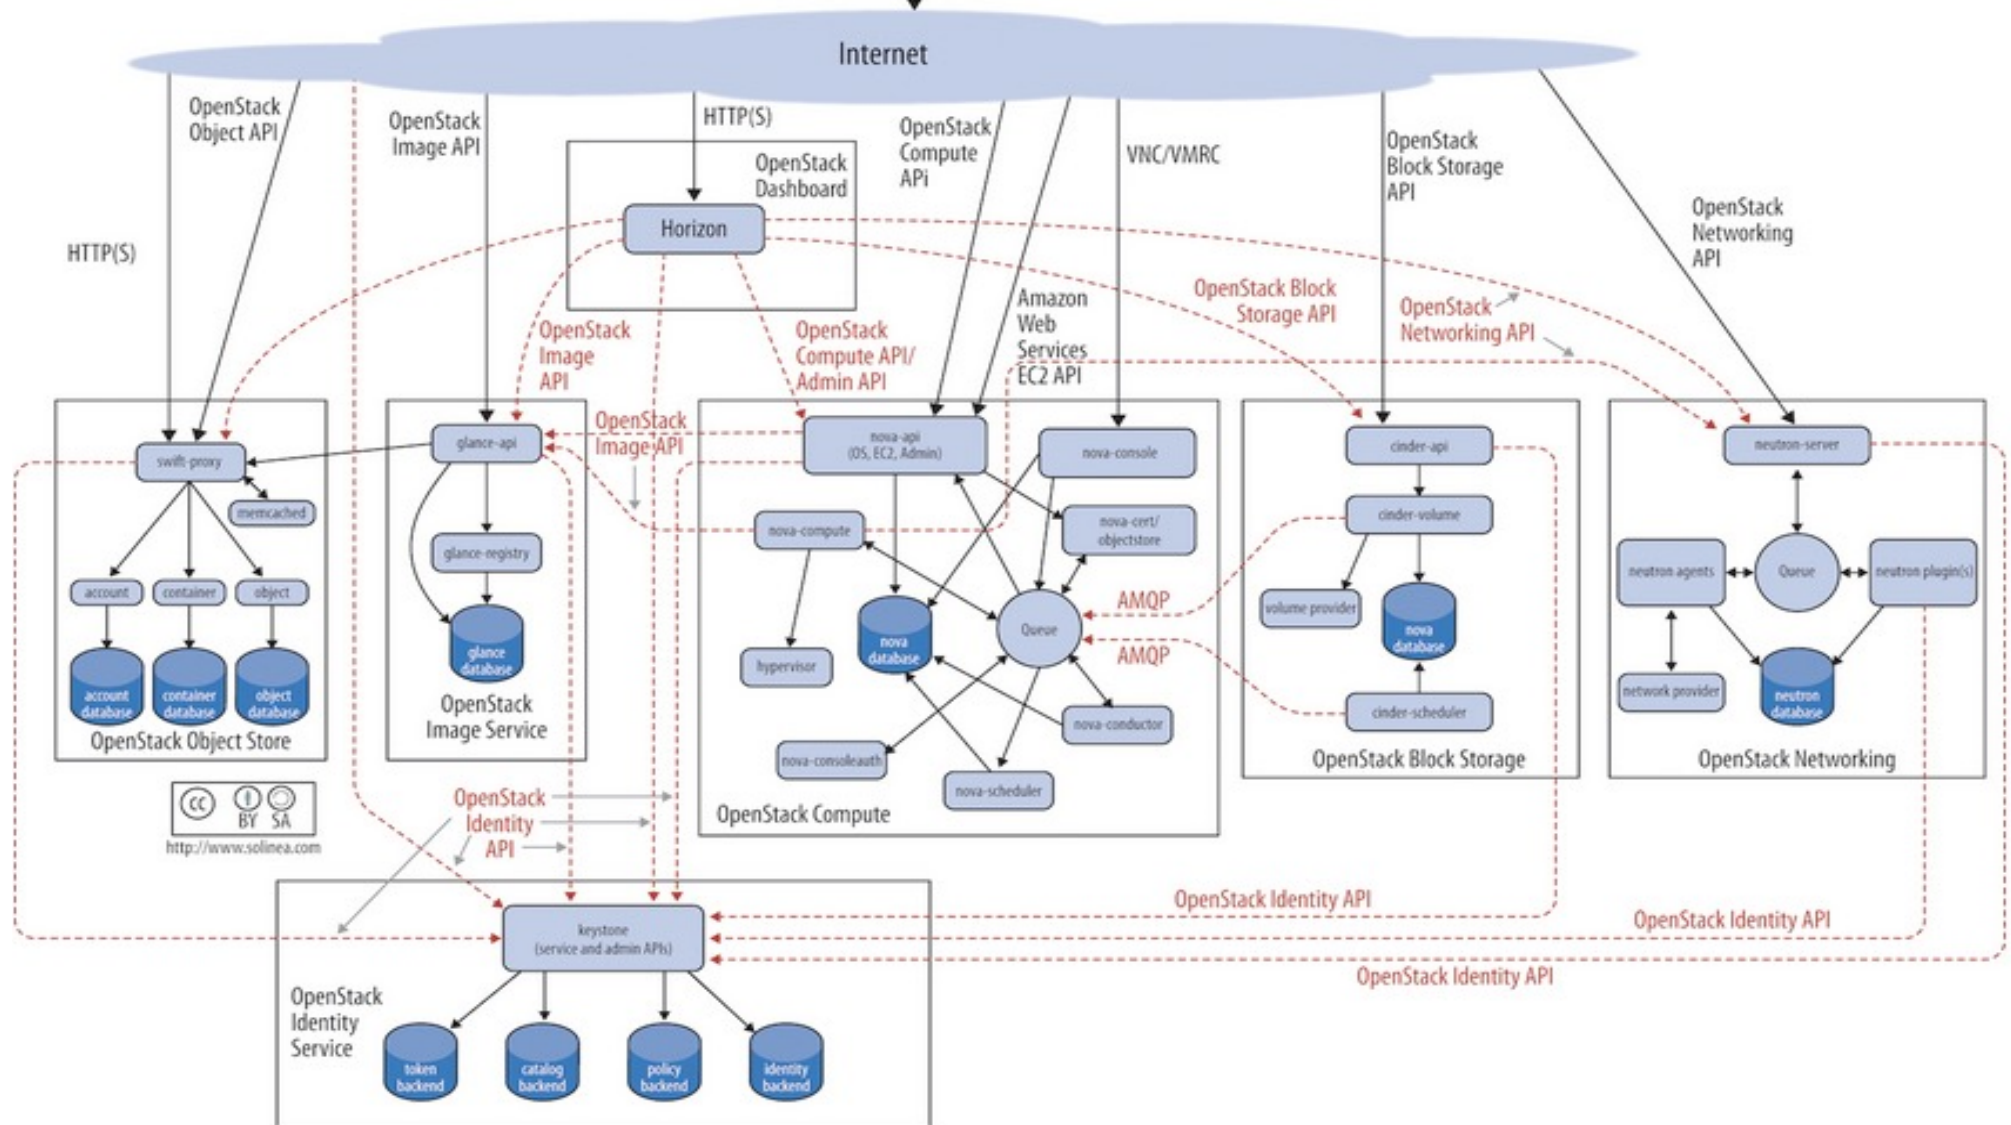
<https://www.openstack.org/software/project-navigator/openstack-components#openstack-services>

(Simplified) OpenStack Architecture

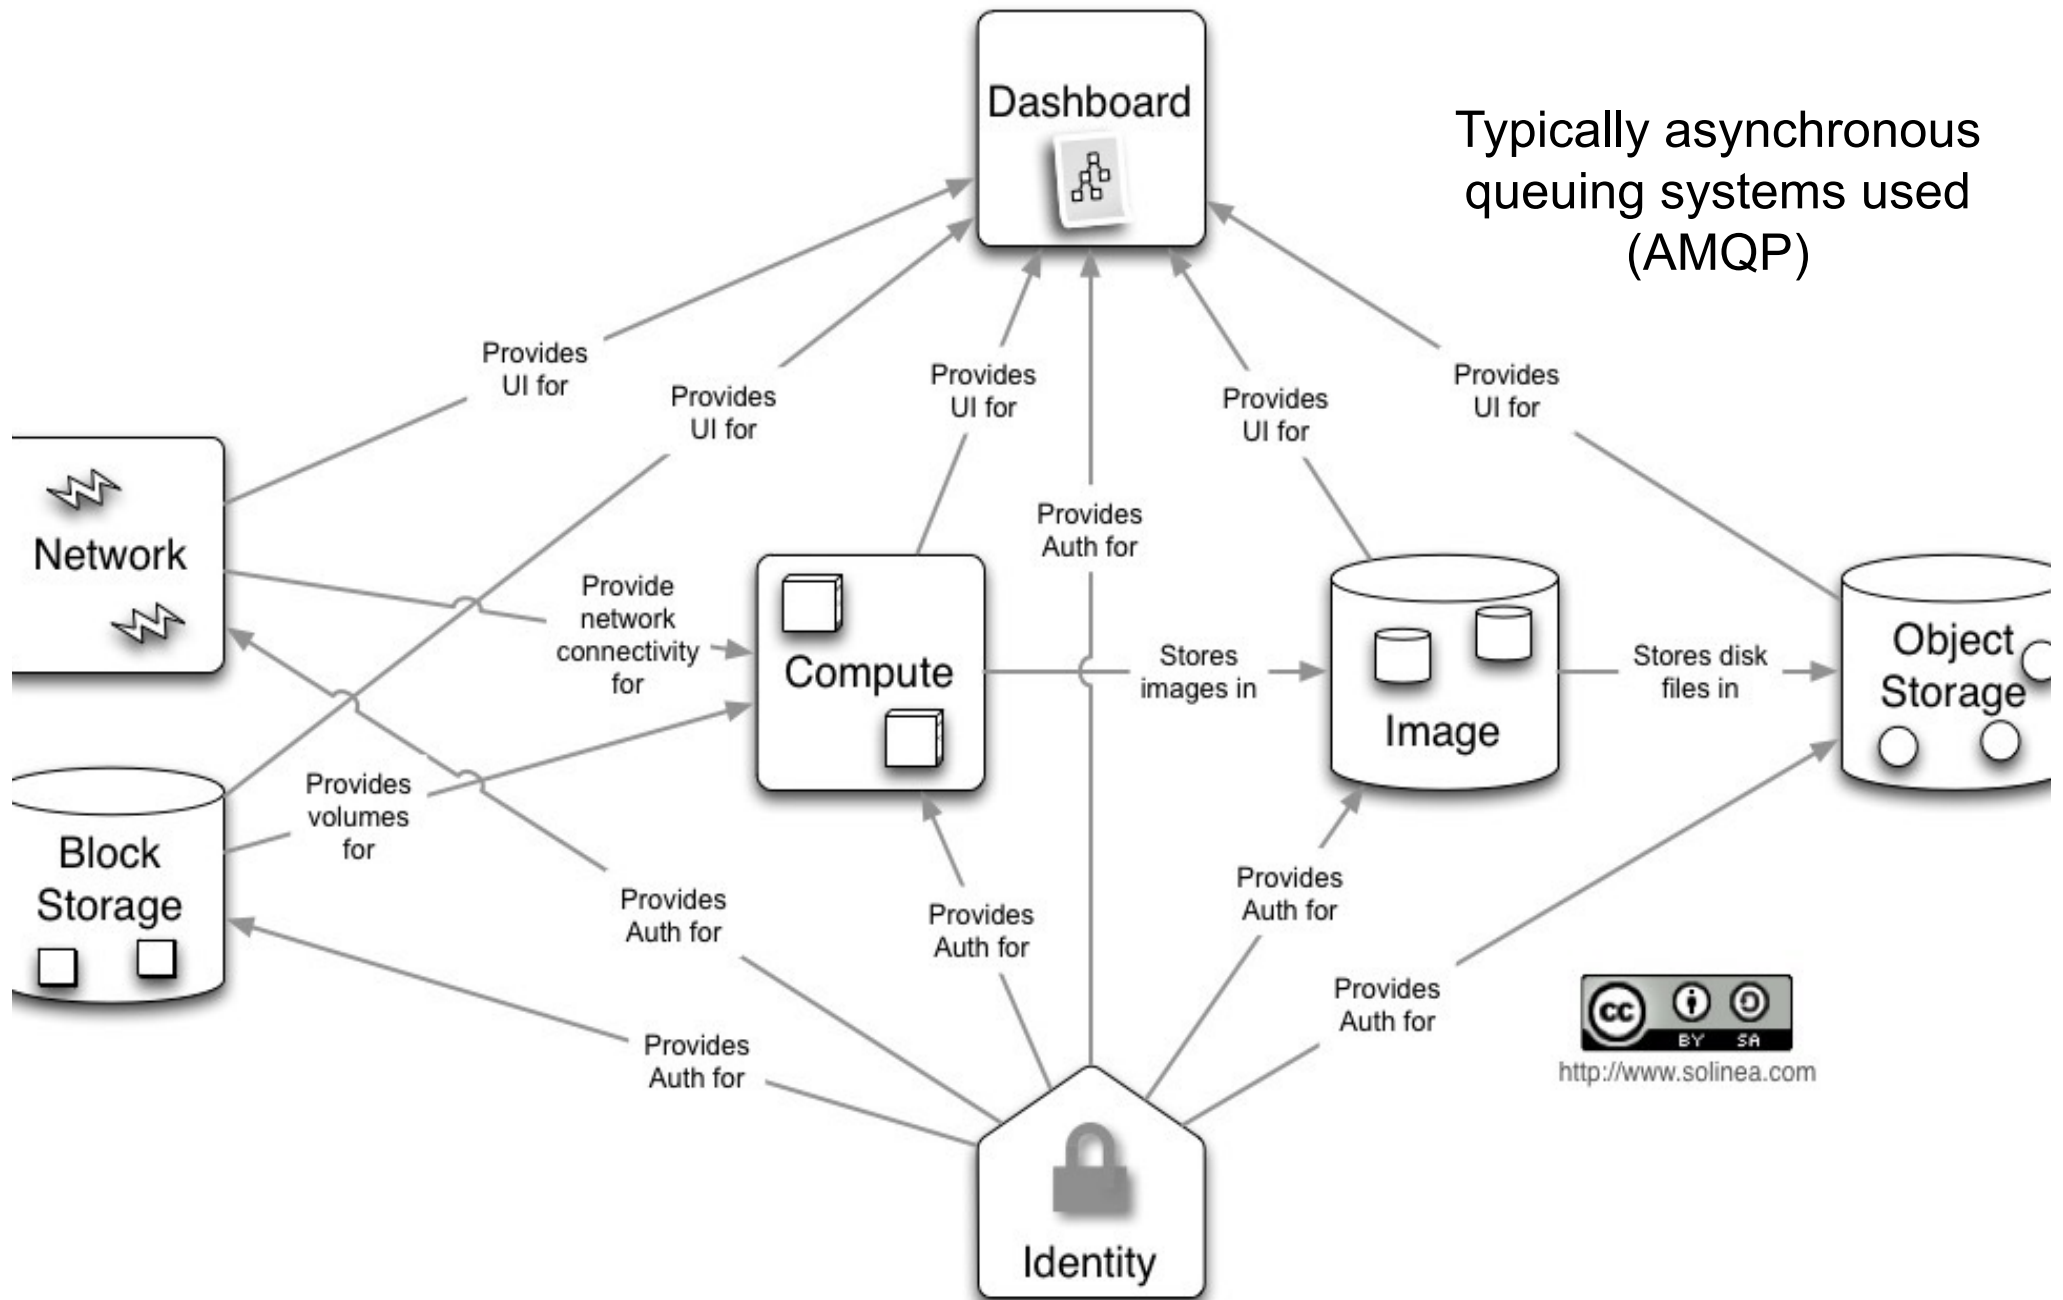


- Command-line interfaces (nova, neutron, swift, etc)
- Cloud Management Tools (Rightscale, Enstratus, etc)
- GUI tools (Dashboard, Cyberduck, iPhone client, etc)

*Components realised as
APIs, Services, Daemons,
Clients, ...*



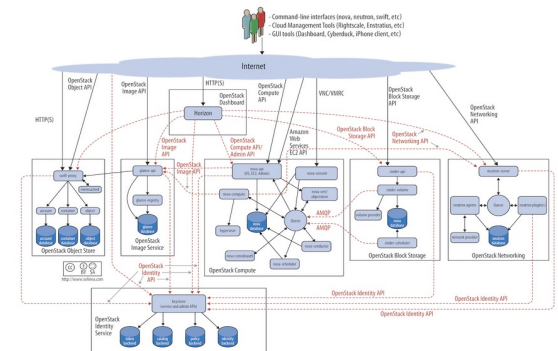
(Simplified) User Perspective



Key Services::Identity Service

- Keystone

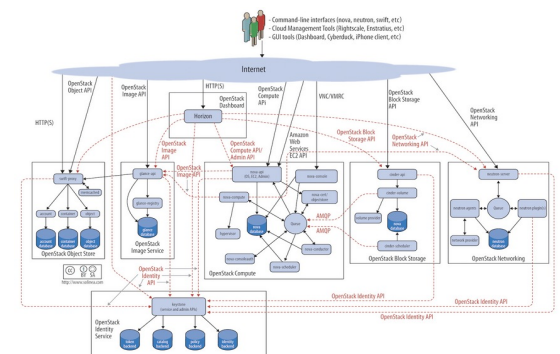
- Provides an authentication and authorization* service for OpenStack services
 - Tracks users/permissions
- Provides a catalog of endpoints for all OpenStack services
 - Each service registered during install
 - Know where they are and who can do what with them
 - Project membership; firewall rules; image mgt; ...
- *Generic authorization system for openStack...
 - more in security lecture



Key Services::Compute

- Nova

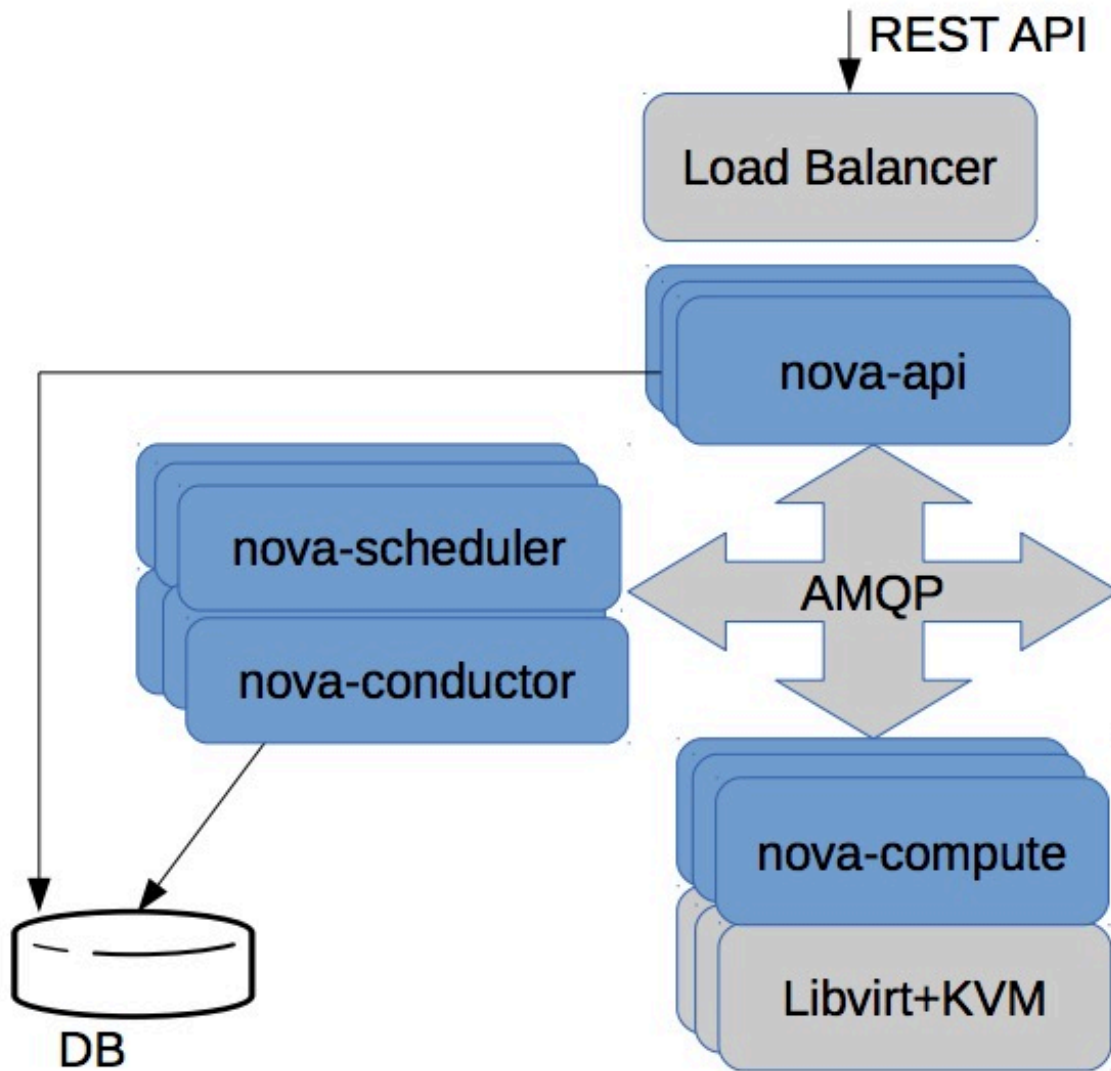
- Manages the lifecycle of compute instances in an OpenStack environment
- Responsibilities include spawning, scheduling and decommissioning of virtual machines on demand
- Virtualisation agnostic
 - Libvirt
 - open source API, daemon and tools for managing platform virtualisation including support for Kernel based virtual machine (KVM), Quick Emulator (QEMU), Xen, Lightweight Linux Container System (LXC)
 - XenAPI, Hyper-V, VMWare ESX,
 - Docker (more later from Luca)
 - ...



Key Services::Compute

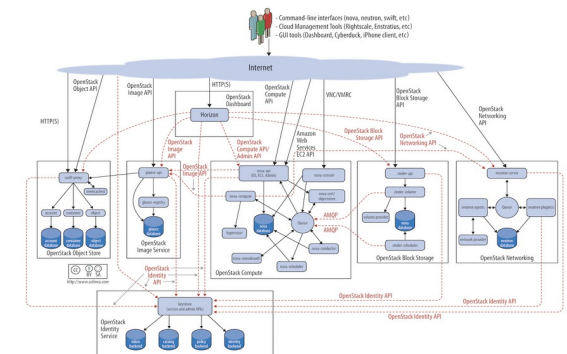
- Nova
 - API
 - **Nova-api** - accepts/responds to end user API calls; supports openStack Compute & EC2 & admin APIs
 - Compute Core
 - **Nova-compute** - Daemon that creates/terminates VMs through hypervisor APIs
 - **Nova-scheduler** - schedules VM instance requests from queue and determines which server host to run
 - **Nova-conductor** - Mediates interactions between compute services and other components, e.g. image database
 - Networking
 - **Nova-network** - Accepts network tasks from queue and manipulates network, e.g. changing IPtable rules
 - Image Mgt, Client Tools, ...

Simplified (Scalable) Nova Architecture



I need a VM with:

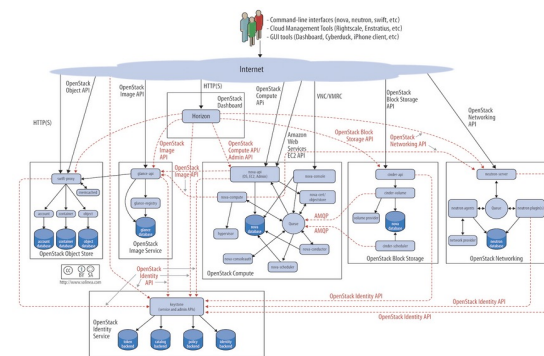
- 64Gb memory,
- 8vCPUs,
- in Melbourne,
- running Ubuntu 12.04,
- ...



Key Services::Object Storage

- **Swift**

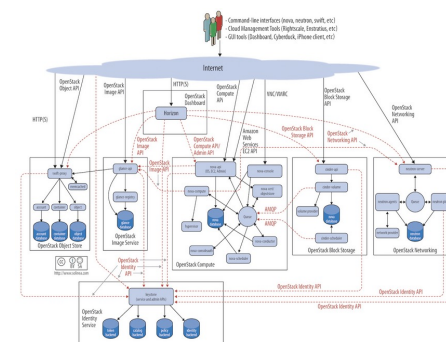
- Stores and retrieves arbitrary unstructured data objects via RESTful API, e.g. VM images and data
 - Not POSIX (atomic operations); eventual consistency
- Fault tolerant with data replication and scale-out architecture.
 - Available from anywhere; persists until deleted
 - Allows to write objects and files to multiple drives, ensuring the data is replicated across a server cluster
- Can be used with/without Nova/compute
- Client; admin support
 - e.g. **Swift client** – allows users to submit commands to ReST API through command line clients to configure/connect object storage to VMs



Key Services::Block Storage

- Cinder

- Provides persistent block storage to virtual machines (instances) and supports creation and management of block storage devices
- Cinder access associated with a VM
 - **Cinder-api** – routes requests to cinder-volume
 - **Cinder-volume** – interacts with block storage service and scheduler to read/write requests; can interact with multiple flavours of storage (flexible driver architecture)
 - **Cinder-scheduler** – selects optimal storage provider node to create volumes (ala nova-scheduler)
 - **Cinder-backup** – provides backup to any types of volume to backup storage provider
 - Can interact with variety of storage solutions



Key Services::Image Service

- Glance


- Accepts requests for disk or server images and their associated metadata (from Swift) and retrieves / installs (through Nova)

- Glance-api – image discovery, retrieval and storage requests

- Glance-registry – stores, processes and retrieves metadata about images, e.g. size and type

- Ubuntu 14.04...?

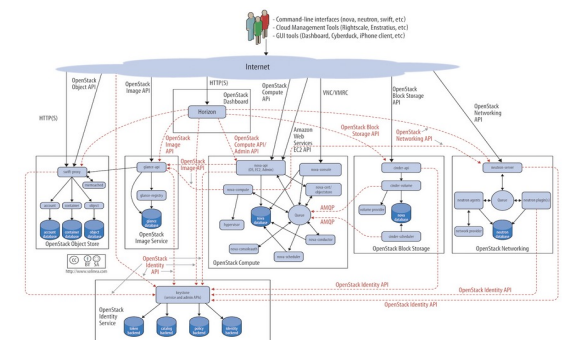
- My last good snapshot...?



Key Services::Networking

- Neutron

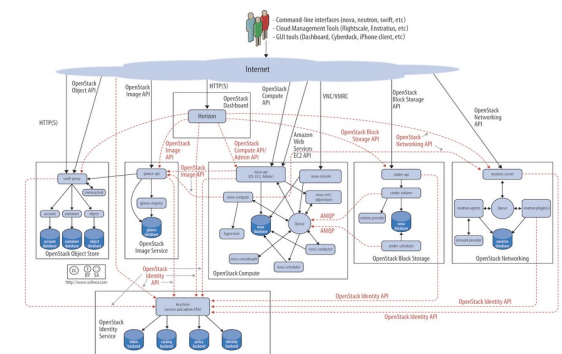
- Supports networking of OpenStack services
- Offers an API for users to define networks and the attachments into them, e.g. switches, routers
- Pluggable architecture that supports multiple networking vendors and technologies
- **Neutron-server** – accepts and routes API requests to appropriate plug-ins for action
 - Port management, e.g. default SSH, VM-specific rules, ...
 - More broadly configuration of availability zone networking, e.g. subnets, DHCP, ...



Key Services::Dashboard

- Horizon

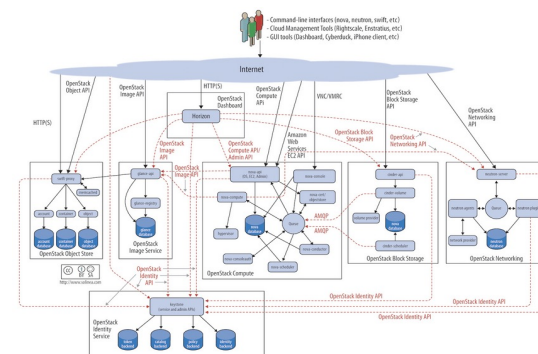
- Provides a web-based self-service portal to interact with underlying OpenStack services, such as launching an instance, assigning IP addresses and configuring access controls.
- Based on Python/Django web application
- Mod_wsgi
 - Apache plug realising web service gateway interface
- Requires Nova, Keystone, Glance, Neutron
- Other services optional...



Key Services::Database Service

- Trove

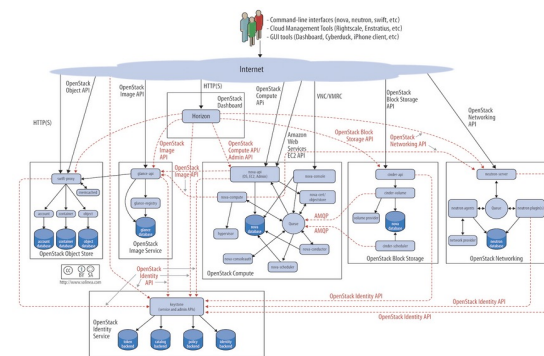
- Provides scalable and reliable Cloud database (DBaaS) functionality for both relational and non-relational database engines (for the masses!)
 - Resource isolation, high performance, automates deployment, config, patching, backups, restores, monitoring...
 - e.g. Set up 3 VMs with mySQL, CouchDB, MongoDB
 - Use image service for each DB type and **trove-manage** to offer them to tenants/user communities



Key Services::Data Processing Service

- Sahara

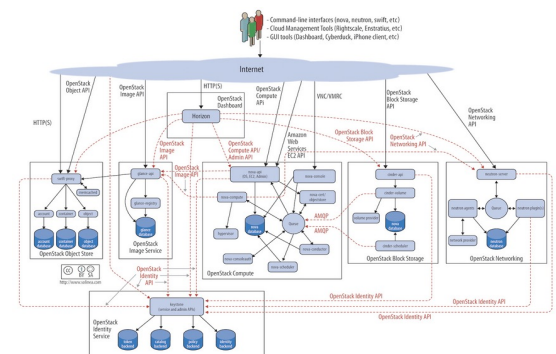

- Provides capabilities to provision and scale Hadoop clusters in OpenStack by specifying parameters such as Hadoop version, cluster topology and node hardware details
 - User fills in details and Sahara supports the automated deployment of infrastructure with support for addition/removal of worker nodes on demand



Key Services::Orchestration Service

- Heat

- Template-driven service to manage lifecycle of applications deployed on Openstack
- *Stack*: Another name for the template and procedure behind creating infrastructure and the required resources from the template file
- Can be integrated with automation tools such as Chef, Puppet, Ansible, etc.
- Template format:



Key Services::Orchestration Service

- Heat details

- *heat_template_version*: allows to specify which version of Heat, the template was written for (*optional*)
- *Description*: describes the intent of the template to a human audience (*optional*)
- *Parameters*: the arguments that the user might be required to provide (*optional*)
- *Resources*: the specifications of resources that are to be created (*mandatory*)
- *Outputs*: any expected values that are to be returned once the template has been processed (*optional*)

Creating Stacks in MRC/NeCTAR

- 1) Create the template file according to your requirements
- 2) Provide environment details (name of key file, image id, etc)
- 3) Select a name for your stack and confirm the parameters
- 4) Make sure rollback checkbox is marked, so if anything goes wrong, all partially created resources get dumped too
- 5) Wait for the magic to happen!

Demonstration of HEAT

- Creating a Wordpress website MRC/NeCTAR-style
- Creating a Wordpress website AWS-style

References

- 1) NeCTAR sample template repository
(<https://github.com/NeCTAR-RC/heat-templates>)