

AI Planning for Autonomy

2. Search Algorithms

Basic Stuff You're Gonna Need to Search for a Solution
Where To Search Next?

Chris Ewin & Tim Miller



THE UNIVERSITY OF
MELBOURNE

With slides by Nir Lipovetsky

Basic State Model: Classical Planning

Ambition:

Write one program that can solve all classical search problems.

State Model $\mathcal{S}(P)$:

- finite and discrete state space S
- a **known initial state** $s_0 \in S$
- a set $S_G \subseteq S$ of goal states
- actions $A(s) \subseteq A$ applicable in each $s \in S$
- a **deterministic transition function** $s' = f(a, s)$ for $a \in A(s)$
- positive **action costs** $c(a, s)$

→ A **solution** is a sequence of applicable actions that maps s_0 into S_G , and it is **optimal** if it minimizes **sum of action costs** (e.g., # of steps)

→ Different **models** and **controllers** obtained by relaxing assumptions in **blue** ...

Example

Criteria for Evaluating Search Strategies

Guarantees:

Completeness: Is the strategy guaranteed to find a solution when there is one?

Optimality: Are the returned solutions guaranteed to be optimal?

Criteria for Evaluating Search Strategies

Guarantees:

Completeness: Is the strategy guaranteed to find a solution when there is one?

Optimality: Are the returned solutions guaranteed to be optimal?

Complexity:

Time Complexity: How long does it take to find a solution? (Measured in **generated states**.)

Space Complexity: How much memory does the search require? (Measured in **states**.)

Criteria for Evaluating Search Strategies

Guarantees:

Completeness: Is the strategy guaranteed to find a solution when there is one?

Optimality: Are the returned solutions guaranteed to be optimal?

Complexity:

Time Complexity: How long does it take to find a solution? (Measured in **generated states**.)

Space Complexity: How much memory does the search require? (Measured in **states**.)

Typical state space features governing complexity:

Branching factor b : How many successors does each state have?

Goal depth d : The number of actions required to reach the shallowest goal state.

Agenda


1 Blind Systematic Search Algorithms

2 Informed Systematic Search Algorithms

Breadth-First Search: Illustration and Guarantees

Strategy: Expand nodes in the order they were produced (FIFO frontier).

Illustration:



Breadth-First Search: Illustration and Guarantees

Strategy: Expand nodes in the order they were produced (FIFO frontier).


Illustration:



Breadth-First Search: Illustration and Guarantees

Strategy: Expand nodes in the order they were produced (FIFO frontier).


Illustration:



Breadth-First Search: Illustration and Guarantees

Strategy: Expand nodes in the order they were produced (FIFO frontier).


Illustration:



Breadth-First Search: Illustration and Guarantees

Strategy: Expand nodes in the order they were produced (FIFO frontier).

Illustration:




Guarantees:

- Completeness?

Breadth-First Search: Illustration and Guarantees

Strategy: Expand nodes in the order they were produced (FIFO frontier).

Illustration:




Guarantees:

- **Completeness?** Yes.
- **Optimality?**

Breadth-First Search: Illustration and Guarantees

Strategy: Expand nodes in the order they were produced (FIFO frontier).

Illustration:



Guarantees:

- **Completeness?** Yes.
- **Optimality?** Yes, for uniform action costs. Breadth-first search always finds a shallowest goal state. If costs are not uniform, this is not necessarily optimal.

Breadth-First Search: Complexity

Time Complexity: Say that b is the maximal branching factor, and d is the goal depth (depth of shallowest goal state).

- Upper bound on the number of generated nodes?

Breadth-First Search: Complexity

Time Complexity: Say that b is the maximal branching factor, and d is the goal depth (depth of shallowest goal state).

- **Upper bound on the number of generated nodes?** $b + b^2 + b^3 + \dots + b^d$: In the worst case, the algorithm generates all nodes in the first d layers.
- So the time complexity is $O(b^d)$.
- **And if we were to apply the goal test at node-expansion time, rather than node-generation time?**

Breadth-First Search: Complexity

Time Complexity: Say that b is the maximal branching factor, and d is the goal depth (depth of shallowest goal state).

- **Upper bound on the number of generated nodes?** $b + b^2 + b^3 + \dots + b^d$: In the worst case, the algorithm generates all nodes in the first d layers.
- So the time complexity is $O(b^d)$.
- **And if we were to apply the goal test at node-expansion time, rather than node-generation time?** $O(b^{d+1})$ because then we'd generate the first $d + 1$ layers in the worst case.

Space Complexity:

Breadth-First Search: Complexity

Time Complexity: Say that b is the maximal branching factor, and d is the goal depth (depth of shallowest goal state).

- **Upper bound on the number of generated nodes?** $b + b^2 + b^3 + \dots + b^d$: In the worst case, the algorithm generates all nodes in the first d layers.
- So the time complexity is $O(b^d)$.
- **And if we were to apply the goal test at node-expansion time, rather than node-generation time?** $O(b^{d+1})$ because then we'd generate the first $d + 1$ layers in the worst case.

Space Complexity: Same as time complexity since all generated nodes are kept in memory.

Breadth-First Search: Example Data

Setting: $b = 10$; 10000 nodes/second; 1000 bytes/node.

Yields data: (inserting values into previous equations)

Depth	Nodes	Time		Memory	
2	110	.11	milliseconds	107	kilobytes
4	11110	11	milliseconds	10.6	megabytes
6	10^6	1.1	seconds	1	gigabyte
8	10^8	2	minutes	103	gigabytes
10	10^{10}	3	hours	10	terabytes
12	10^{12}	13	days	1	petabyte
14	10^{14}	3.5	years	99	petabytes

→ So, which is the worse problem, time or memory?


Breadth-First Search: Example Data

Setting: $b = 10$; 10000 nodes/second; 1000 bytes/node.

Yields data: (inserting values into previous equations)

Depth	Nodes	Time		Memory	
2	110	.11	milliseconds	107	kilobytes
4	11110	11	milliseconds	10.6	megabytes
6	10^6	1.1	seconds	1	gigabyte
8	10^8	2	minutes	103	gigabytes
10	10^{10}	3	hours	10	terabytes
12	10^{12}	13	days	1	petabyte
14	10^{14}	3.5	years	99	petabytes


→ So, which is the worse problem, time or memory? Memory. (In my own experience, typically exhausts RAM memory within a few minutes.)




Depth-First Search: Illustration

Strategy: Expand the most recent nodes in (LIFO frontier).

Illustration: (Nodes at depth 3 are assumed to have no successors)






Depth-First Search: Illustration

Strategy: Expand the most recent nodes in (LIFO frontier).

Illustration: (Nodes at depth 3 are assumed to have no successors)



Depth-First Search: Guarantees and Complexity

Guarantees:

Question: A) Complete and optimal B) Complete but may not be optimal C) Optimal but may not be complete D) Neither complete nor optimal

Depth-First Search: Guarantees and Complexity

Guarantees:

Question: A) Complete and optimal B) Complete but may not be optimal C) Optimal but may not be complete D) Neither complete nor optimal

- Optimality?

Depth-First Search: Guarantees and Complexity

Guarantees:

Question: A) Complete and optimal B) Complete but may not be optimal C) Optimal but may not be complete D) Neither complete nor optimal

- **Optimality?** No. After all, the algorithm just “chooses some direction and hopes for the best”. (Depth-first search is a way of “hoping to get lucky”.)
- **Completeness?**

Depth-First Search: Guarantees and Complexity

Guarantees:

Question: A) Complete and optimal B) Complete but may not be optimal C) Optimal but may not be complete D) Neither complete nor optimal

- **Optimality?** No. After all, the algorithm just “chooses some direction and hopes for the best”. (Depth-first search is a way of “hoping to get lucky”.)
- **Completeness?** No, because search branches may be infinitely long: No check for cycles along a branch!
→ Depth-first search is complete in case the state space is **acyclic**, e.g., **Constraint Satisfaction Problems**. If we do add a cycle check, it becomes complete for finite state spaces.

Depth-First Search: Guarantees and Complexity

Guarantees:

Question: A) Complete and optimal B) Complete but may not be optimal C) Optimal but may not be complete D) Neither complete nor optimal

- **Optimality?** No. After all, the algorithm just “chooses some direction and hopes for the best”. (Depth-first search is a way of “hoping to get lucky”.)
- **Completeness?** No, because search branches may be infinitely long: No check for cycles along a branch!
→ Depth-first search is complete in case the state space is **acyclic**, e.g., **Constraint Satisfaction Problems**. If we do add a cycle check, it becomes complete for finite state spaces.

Complexity:

- **Space:** Stores nodes and applicable actions on the path to the current node. So if m is the maximal depth reached, the complexity is

Depth-First Search: Guarantees and Complexity

Guarantees:

Question: A) Complete and optimal B) Complete but may not be optimal C) Optimal but may not be complete D) Neither complete nor optimal

- **Optimality?** No. After all, the algorithm just “chooses some direction and hopes for the best”. (Depth-first search is a way of “hoping to get lucky”.)
- **Completeness?** No, because search branches may be infinitely long: No check for cycles along a branch!
→ Depth-first search is complete in case the state space is **acyclic**, e.g., **Constraint Satisfaction Problems**. If we do add a cycle check, it becomes complete for finite state spaces.

Complexity:

- **Space:** Stores nodes and applicable actions on the path to the current node. So if m is the maximal depth reached, the complexity is $O(bm)$.
- **Time:** If there are paths of length m in the state space, $O(b^m)$ nodes can be generated. Even if there are solutions of depth 1!

Depth-First Search: Guarantees and Complexity

Guarantees:

Question: A) Complete and optimal B) Complete but may not be optimal C) Optimal but may not be complete D) Neither complete nor optimal

- **Optimality?** No. After all, the algorithm just “chooses some direction and hopes for the best”. (Depth-first search is a way of “hoping to get lucky”.)
- **Completeness?** No, because search branches may be infinitely long: No check for cycles along a branch!
→ Depth-first search is complete in case the state space is **acyclic**, e.g., **Constraint Satisfaction Problems**. If we do add a cycle check, it becomes complete for finite state spaces.

Complexity:


- **Space:** Stores nodes and applicable actions on the path to the current node. So if m is the maximal depth reached, the complexity is $O(bm)$.
- **Time:** If there are paths of length m in the state space, $O(b^m)$ nodes can be generated. Even if there are solutions of depth 1!
→ If we happen to choose “the right direction” then we can find a length- l solution in time $O(bl)$ regardless how big the state space is.

Iterative Deepening Search: Illustration


Limit = 0



Iterative Deepening Search: Illustration



Iterative Deepening Search: Illustration



Iterative Deepening Search: Illustration






Limit = 0





Limit = 1



Limit = 2



Limit = 3



Iterative Deepening Search: Guarantees and Complexity

*“Iterative Deepening Search=
Keep doing the same work over again until you find a solution.”*

Iterative Deepening Search: Guarantees and Complexity

*“Iterative Deepening Search=
Keep doing the same work over again until you find a solution.”*

BUT: Optimality?

Iterative Deepening Search: Guarantees and Complexity

*“Iterative Deepening Search=
Keep doing the same work over again until you find a solution.”*

BUT: Optimality? Yes! (assuming uniform costs) Completeness?

Iterative Deepening Search: Guarantees and Complexity

*“Iterative Deepening Search=
Keep doing the same work over again until you find a solution.”*

BUT: Optimality? Yes! (assuming uniform costs) Completeness? Yes! Space complexity?

Iterative Deepening Search: Guarantees and Complexity

*“Iterative Deepening Search=
Keep doing the same work over again until you find a solution.”*

BUT: **Optimality?** Yes! (assuming uniform costs) **Completeness?** Yes! **Space complexity?** $O(b d)$.

Iterative Deepening Search: Guarantees and Complexity

*“Iterative Deepening Search=
Keep doing the same work over again until you find a solution.”*

BUT: **Optimality?** Yes! (assuming uniform costs) **Completeness?** Yes! **Space complexity?** $O(bd)$.

Time complexity:

Breadth-First-Search	$b + b^2 + \dots + b^{d-1} + b^d \in O(b^d)$
Iterative Deepening Search	$(d)b + (d-1)b^2 + \dots + 3b^{d-2} + 2b^{d-1} + 1b^d \in O(b^d)$

Iterative Deepening Search: Guarantees and Complexity

*“Iterative Deepening Search=
Keep doing the same work over again until you find a solution.”*

BUT: **Optimality?** Yes! (assuming uniform costs) **Completeness?** Yes! **Space complexity?** $O(bd)$.

Time complexity:

Breadth-First-Search	$b + b^2 + \dots + b^{d-1} + b^d \in O(b^d)$
Iterative Deepening Search	$(d)b + (d-1)b^2 + \dots + 3b^{d-2} + 2b^{d-1} + 1b^d \in O(b^d)$

Example: $b = 10, d = 5$

Breadth-First Search	$10 + 100 + 1,000 + 10,000 + 100,000 = 111,110$
Iterative Deepening Search	$50 + 400 + 3,000 + 20,000 + 100,000 = 123,450$

→ IDS combines the advantages of breadth-first and depth-first search. It is the preferred blind search method in large state spaces with unknown solution depth.

Agenda

1 Blind Systematic Search Algorithms

2 Informed Systematic Search Algorithms

Greedy Best-First Search

Greedy Best-First Search (with duplicate detection)

```
open := new priority queue ordered by ascending  $h(\text{state}(\sigma))$ 
open.insert(make-root-node(init()))
closed :=  $\emptyset$ 
while not open.empty():
     $\sigma := \text{open.pop-min()}$  /* get best state */
    if  $\text{state}(\sigma) \notin \text{closed}$ : /* check duplicates */
        closed := closed  $\cup \{\text{state}(\sigma)\}$  /* close state */
        if is-goal(state( $\sigma$ )): return extract-solution( $\sigma$ )
        for each  $(a, s') \in \text{succ}(\text{state}(\sigma))$ : /* expand state */
             $\sigma' := \text{make-node}(\sigma, a, s')$ 
            if  $h(\text{state}(\sigma')) < \infty$ : open.insert( $\sigma'$ )
return unsolvable
```

Greedy Best-First Search: Remarks

Properties:

- Complete?

¹Even for perfect heuristics! E.g., say the start state has two transitions to goal states, one of which costs a million bucks while the other one is free. Nothing keeps Greedy Best-First Search from choosing the bad one.

Greedy Best-First Search: Remarks

Properties:

- **Complete?** Yes, for safe heuristics. (and duplicate detection to avoid cycles)
- **Optimal?**

¹ Even for perfect heuristics! E.g., say the start state has two transitions to goal states, one of which costs a million bucks while the other one is free. Nothing keeps Greedy Best-First Search from choosing the bad one.

Greedy Best-First Search: Remarks

Properties:

- **Complete?** Yes, for safe heuristics. (and duplicate detection to avoid cycles)
- **Optimal?** No.¹
- Invariant under all strictly monotonic transformations of h (e.g., scaling with a positive constant or adding a constant).

Implementation:

- Priority queue: e.g., a [min heap](#).
- “Check Duplicates”: Could already do in “expand state”; done here after “get best state” *only* to more clearly point out relation to A*.

¹ Even for perfect heuristics! E.g., say the start state has two transitions to goal states, one of which costs a million bucks while the other one is free. Nothing keeps Greedy Best-First Search from choosing the bad one.

A*

A* (with duplicate detection and re-opening)

```

open := new priority queue ordered by ascending  $g(\text{state}(\sigma)) + h(\text{state}(\sigma))$ 
open.insert(make-root-node(init()))
closed :=  $\emptyset$ 
best-g :=  $\emptyset$  /* maps states to numbers */
while not open.empty():
     $\sigma := \text{open.pop-min}()$ 
    if  $\text{state}(\sigma) \notin \text{closed}$  or  $g(\sigma) < \text{best-g}(\text{state}(\sigma))$ :
        /* re-open if better g; note that all  $\sigma'$  with same state but worse g
           are behind  $\sigma$  in open, and will be skipped when their turn comes */
        closed := closed  $\cup$  {state( $\sigma$ )}
        best-g(state( $\sigma$ )) :=  $g(\sigma)$ 
        if is-goal(state( $\sigma$ )): return extract-solution( $\sigma$ )
        for each  $(a, s') \in \text{succ}(\text{state}(\sigma))$ :
             $\sigma' := \text{make-node}(\sigma, a, s')$ 
            if  $h(\text{state}(\sigma')) < \infty$ : open.insert( $\sigma'$ )
return unsolvable

```


A*: Remarks

Properties:

- Complete?

A*: Remarks

Properties:

- **Complete?** Yes, for safe heuristics. (Even without duplicate detection.)
- **Optimal?**

A*: Remarks

Properties:

- **Complete?** Yes, for safe heuristics. (Even without duplicate detection.)
- **Optimal?** Yes, for admissible heuristics. (Even without duplicate detection.)

Implementation:

- Popular method: break ties ($f(s) = f(s')$) by smaller h -value.
- If h is admissible and consistent, then A* never re-opens a state. So if we know that this is the case, then we can simplify the algorithm.
- Common, hard to spot bug: check duplicates at the wrong point. (Russel & Norvig are way too imprecise about this.)
- Our implementation is optimized for readability not for efficiency!

Question

Question!

If we set $h(n) := 0$ for all n , what does A^* become?

- (A): Breadth-first search.
- (B): Depth-first search.
- (C): Uniform-cost search.
- (D): Depth-limited search.

Question

Question!

If we set $h(n) := 0$ for all n , what does A^* become?

- | | |
|----------------------------|----------------------------|
| (A): Breadth-first search. | (B): Depth-first search. |
| (C): Uniform-cost search. | (D): Depth-limited search. |

→ (C): Same expansion order. (Details in book-keeping of open/closed states may differ.)