

AI Planning for Autonomy

2. Search Algorithms

Basic Stuff You're Gonna Need to Search for a Solution
Where To Search Next?

Chris Ewin & Tim Miller



THE UNIVERSITY OF
MELBOURNE

With slides by Nir Lipovetsky

Basic State Model: Classical Planning

Ambition:

Write one program that can solve all classical search problems.

State Model $\mathcal{S}(P)$:

- finite and discrete state space S
- a **known initial state** $s_0 \in S$
- a set $S_G \subseteq S$ of goal states
- actions $A(s) \subseteq A$ applicable in each $s \in S$
- a **deterministic transition function** $s' = f(a, s)$ for $a \in A(s)$
- positive **action costs** $c(a, s)$

→ A **solution** is a sequence of applicable actions that maps s_0 into S_G , and it is **optimal** if it minimizes **sum of action costs** (e.g., # of steps)

→ Different **models** and **controllers** obtained by relaxing assumptions in **blue** ...

Solving the State Model: Path-finding in graphs

Search algorithms for planning exploit the **correspondence** between **(classical) states model** $\mathcal{S}(P)$ and **directed graphs**:

- The **nodes** of the graph represent the **states** s in the model
- The edges (s, s') capture corresponding transition in the model with same cost

In the **planning as heuristic search** formulation, the problem P is solved by **path-finding** algorithms over the **graph** associated with model $\mathcal{S}(P)$

Classification of Search Algorithms

Blind search vs. heuristic (or informed) search:

- **Blind search algorithms:** Only use the basic ingredients for general search algorithms.
 - *e.g., Depth First Search (DFS), Breadth-first search (BrFS), Uniform Cost (Dijkstra), Iterative Deepening (ID)*
- **Heuristic search algorithms:** Additionally use **heuristic functions** which estimate the distance (or remaining cost) to the goal.
 - *e.g., A*, IDA*, Hill Climbing, Best First, WA*, DFS B&B, LRTA*, ...*

Systematic search vs. local search:

- **Systematic search algorithms:** Consider a large number of search nodes simultaneously.
- **Local search algorithms:** Work with one (or a few) candidate solutions (search nodes) at a time.
 - This is not a black-and-white distinction; there are *crossbreeds* (e.g., **enforced hill-climbing**).

What works where in planning?

Blind search vs. heuristic search:

- For **satisficing** planning, heuristic search vastly outperforms blind algorithms pretty much everywhere.
- For **optimal** planning, heuristic search also is better (but the difference is less pronounced).

Systematic search vs. local search:

- For **satisficing** planning, there are successful instances of each.
- For **optimal** planning, systematic algorithms are required.

What works where in planning?

Blind search vs. heuristic search:

- For **satisficing** planning, heuristic search vastly outperforms blind algorithms pretty much everywhere.
- For **optimal** planning, heuristic search also is better (but the difference is less pronounced).

Systematic search vs. local search:

- For **satisficing** planning, there are successful instances of each.
- For **optimal** planning, systematic algorithms are required.

→ Here, we cover the subset of search algorithms most successful in planning. Only some Blind search algorithms are covered. (refer to Russel & Norvig Chapters 3 and 4 for that).

Agenda

- 1 Basics
- 2 Blind Systematic Search Algorithms
- 3 Heuristic Functions
- 4 Informed Systematic Search Algorithms
- 5 Local Search Algorithms
- 6 Conclusion

Search Terminology

Search node n : Contains a *state* reached by the search, plus information about how it was reached.

Search Terminology

Search node n : Contains a *state* reached by the search, plus information about how it was reached.

Path cost $g(n)$: The cost of the path reaching n .

Optimal cost g^* : The cost of an optimal solution path. For a state s , $g^*(s)$ is the cost of a cheapest path reaching s .

Search Terminology

Search node n : Contains a *state* reached by the search, plus information about how it was reached.

Path cost $g(n)$: The cost of the path reaching n .

Optimal cost g^* : The cost of an optimal solution path. For a state s , $g^*(s)$ is the cost of a cheapest path reaching s .

Node expansion: Generating all successors of a node, by applying all actions applicable to the node's state s . Afterwards, the *state* s itself is also said to be expanded.

Search strategy: Method for deciding which node is expanded next.

Search Terminology

Search node n : Contains a *state* reached by the search, plus information about how it was reached.

Path cost $g(n)$: The cost of the path reaching n .

Optimal cost g^* : The cost of an optimal solution path. For a state s , $g^*(s)$ is the cost of a cheapest path reaching s .

Node expansion: Generating all successors of a node, by applying all actions applicable to the node's state s . Afterwards, the *state* s itself is also said to be expanded.

Search strategy: Method for deciding which node is expanded next.

Open list: Set of all *nodes* that currently are candidates for expansion. Also called **frontier**.

Closed list: Set of all *states* that were already expanded. Used only in **graph search**, not in **tree search** (up next). Also called **explored set**.

World States vs. Search States

Reminder: Search Space for Classical Search

A (classical) **search space** is defined by the following three operations:

- **start()**: Generate the start (search) state.
- **is-target(s)**: Test whether a given search state is a target state.
- **succ(s)**: Generates the successor states (a, s') of search state s , along with the actions through which they are reached.

World States vs. Search States

Reminder: Search Space for Classical Search

A (classical) **search space** is defined by the following three operations:

- **start()**: Generate the start (search) state.
- **is-target(s)**: Test whether a given search state is a target state.
- **succ(s)**: Generates the successor states (a, s') of search state s , along with the actions through which they are reached.

Search states \neq world states:

- **Progression:**

World States vs. Search States

Reminder: Search Space for Classical Search

A (classical) **search space** is defined by the following three operations:

- **start()**: Generate the start (search) state.
- **is-target(s)**: Test whether a given search state is a target state.
- **succ(s)**: Generates the successor states (a, s') of search state s , along with the actions through which they are reached.

Search states \neq world states:

- **Progression**: Yes, search states = world states.
- **Regression**:

World States vs. Search States

Reminder: Search Space for Classical Search

A (classical) **search space** is defined by the following three operations:

- **start()**: Generate the start (search) state.
- **is-target(s)**: Test whether a given search state is a target state.
- **succ(s)**: Generates the successor states (a, s') of search state s , along with the actions through which they are reached.

Search states \neq world states:

- **Progression**: Yes, search states = world states.
- **Regression**: No, search states = sets of world states, represented as conjunctive sub-goals.

→ We consider progression in the entire course, unless explicitly stated otherwise.
We use “ s ” to denote world/search states interchangeably.

Search States vs. Search Nodes

- **Search states** s : States (vertices) of the search space.
- **Search nodes** σ : Search states, plus information on where/when/how they are encountered during search.

Search States vs. Search Nodes

- **Search states** s : States (vertices) of the search space.
- **Search nodes** σ : Search states, plus information on where/when/how they are encountered during search.

What is in a search node?

Different search algorithms store different information in a search node σ , but typical information includes:

- $state(\sigma)$: Associated search state.
- $parent(\sigma)$: Pointer to search node from which σ is reached.
- $action(\sigma)$: An action leading from $state(parent(\sigma))$ to $state(\sigma)$.
- $g(\sigma)$: Cost of σ (cost of path from the root node to σ).

For the root node, $parent(\sigma)$ and $action(\sigma)$ are undefined.

Criteria for Evaluating Search Strategies

Guarantees:

Completeness: Is the strategy guaranteed to find a solution when there is one?

Optimality: Are the returned solutions guaranteed to be optimal?

Criteria for Evaluating Search Strategies

Guarantees:

Completeness: Is the strategy guaranteed to find a solution when there is one?

Optimality: Are the returned solutions guaranteed to be optimal?

Complexity:

Time Complexity: How long does it take to find a solution? (Measured in **generated states**.)

Space Complexity: How much memory does the search require? (Measured in **states**.)

Criteria for Evaluating Search Strategies

Guarantees:

Completeness: Is the strategy guaranteed to find a solution when there is one?

Optimality: Are the returned solutions guaranteed to be optimal?

Complexity:

Time Complexity: How long does it take to find a solution? (Measured in **generated states**.)

Space Complexity: How much memory does the search require? (Measured in **states**.)

Typical state space features governing complexity:

Branching factor b : How many successors does each state have?

Goal depth d : The number of actions required to reach the shallowest goal state.

Agenda

- 1 Basics
- 2 Blind Systematic Search Algorithms
- 3 Heuristic Functions
- 4 Informed Systematic Search Algorithms
- 5 Local Search Algorithms
- 6 Conclusion

Before We Begin, ctd.

Blind search strategies we'll discuss:

- **Breadth-first search**. Advantage: time complexity.
Variant: **Uniform cost search**.
- **Depth-first search**. Advantage: space complexity.
- **Iterative deepening search**. Combines advantages of breadth-first search and depth-first search. Uses **depth-limited search** as a sub-procedure.

Before We Begin, ctd.

Blind search strategies we'll discuss:

- **Breadth-first search**. Advantage: time complexity.
Variant: **Uniform cost search**.
- **Depth-first search**. Advantage: space complexity.
- **Iterative deepening search**. Combines advantages of breadth-first search and depth-first search. Uses **depth-limited search** as a sub-procedure.


Blind search strategy we won't discuss:

- **Bi-directional search**. Two separate search spaces, one forward from the initial state, the other backward from the goal. Stops when the two search spaces overlap.

Breadth-First Search: Illustration and Guarantees

Strategy: Expand nodes in the order they were produced (FIFO frontier).

Illustration:



Breadth-First Search: Illustration and Guarantees

Strategy: Expand nodes in the order they were produced (FIFO frontier).


Illustration:



Breadth-First Search: Illustration and Guarantees

Strategy: Expand nodes in the order they were produced (FIFO frontier).


Illustration:



Breadth-First Search: Illustration and Guarantees

Strategy: Expand nodes in the order they were produced (FIFO frontier).


Illustration:



Breadth-First Search: Illustration and Guarantees

Strategy: Expand nodes in the order they were produced (FIFO frontier).

Illustration:




Guarantees:

- Completeness?

Breadth-First Search: Illustration and Guarantees

Strategy: Expand nodes in the order they were produced (FIFO frontier).

Illustration:




Guarantees:

- **Completeness?** Yes.
- **Optimality?**

Breadth-First Search: Illustration and Guarantees

Strategy: Expand nodes in the order they were produced (FIFO frontier).

Illustration:




Guarantees:

- **Completeness?** Yes.
- **Optimality?** Yes, for uniform action costs. Breadth-first search always finds a shallowest goal state. If costs are not uniform, this is not necessarily optimal.

Depth-First Search: Illustration

Strategy: Expand the most recent nodes in (LIFO frontier).


Illustration: (Nodes at depth 3 are assumed to have no successors)



Depth-First Search: Illustration

Strategy: Expand the most recent nodes in (LIFO frontier).


Illustration: (Nodes at depth 3 are assumed to have no successors)



Depth-First Search: Illustration

Strategy: Expand the most recent nodes in (LIFO frontier).


Illustration: (Nodes at depth 3 are assumed to have no successors)



Depth-First Search: Illustration

Strategy: Expand the most recent nodes in (LIFO frontier).

Illustration: (Nodes at depth 3 are assumed to have no successors)



Iterative Deepening Search: Illustration

Limit = 0



Iterative Deepening Search: Illustration

Limit = 0



Limit = 1



Iterative Deepening Search: Illustration



Limit = 0



Limit = 1



Limit = 2



Iterative Deepening Search: Illustration




Limit = 0







Limit = 1



Limit = 2



Limit = 3



Iterative Deepening Search: Guarantees and Complexity

*“Iterative Deepening Search=
Keep doing the same work over again until you find a solution.”*

Agenda

- 1 Basics
- 2 Blind Systematic Search Algorithms
- 3 Heuristic Functions**
- 4 Informed Systematic Search Algorithms
- 5 Local Search Algorithms
- 6 Conclusion

Heuristic Search Algorithms: Systematic

→ Heuristic search algorithms are the most common and overall most successful algorithms for classical planning.

Systematic heuristic search algorithms:

- Greedy best-first search.
 - One of 3 most popular algorithms in satisficing planning.
- Weighted A^* .
 - One of 3 most popular algorithms in satisficing planning.
- A^* .
 - Most popular algorithm in optimal planning. (Rarely ever used for satisficing planning.)
- IDA*, depth-first branch-and-bound search, breadth-first heuristic search, ...


Heuristic Search Algorithms: Local

→ Heuristic search algorithms are the most common and overall most successful algorithms for classical planning.

Local heuristic search algorithms:

- Hill-climbing.
- Enforced hill-climbing.
 - One of 3 most popular algorithms in satisficing planning.
- Beam search, tabu search, genetic algorithms, simulated annealing, ...

Heuristic Search: Basic Idea



→ Heuristic function h estimates the cost of an optimal path to the goal; search gives a preference to explore states with small h .

Heuristic Functions

Heuristic searches require a heuristic function to estimate remaining cost:

Definition (Heuristic Function). Let Π be a planning task with state space Θ_{Π} . A *heuristic function*, short *heuristic*, for Π is a function $h : S \mapsto \mathbb{R}_0^+ \cup \{\infty\}$. Its value $h(s)$ for a state s is referred to as the state's *heuristic value*, or *h-value*.

Definition (Remaining Cost, h^*). Let Π be a planning task with state space Θ_{Π} . For a state $s \in S$, the state's *remaining cost* is the cost of an optimal plan for s , or ∞ if there exists no plan for s . The *perfect heuristic* for Π , written h^* , assigns every $s \in S$ its remaining cost as the heuristic value.

Heuristic Functions: Discussion

What does it mean to “estimate remaining cost”?

- For many heuristic search algorithms, h does not need to have any properties for the algorithm to “work” (= be correct and complete).
→ h is *any* function from states to numbers ...

Heuristic Functions: Discussion

What does it mean to “estimate remaining cost”?

- For many heuristic search algorithms, h does not need to have any properties for the algorithm to “work” (= be correct and complete).
 - h is *any* function from states to numbers . . .
- Search **performance** depends crucially on “how well h reflects h^* ”!!
 - This is informally called the **informedness** or **quality** of h .

Heuristic Functions: Discussion

What does it mean to “estimate remaining cost”?

- For many heuristic search algorithms, h does not need to have any properties for the algorithm to “work” (= be correct and complete).
→ h is *any* function from states to numbers . . .
- Search **performance** depends crucially on “how well h reflects h^* ”!!
→ This is informally called the **informedness** or **quality** of h .
- For some search algorithms, like A^* , we can *prove* relationships between formal quality properties of h and search efficiency (mainly the number of expanded nodes).

Heuristic Functions: Discussion

What does it mean to “estimate remaining cost”?

- For many heuristic search algorithms, h does not need to have any properties for the algorithm to “work” (= be correct and complete).
→ h is *any* function from states to numbers . . .
- Search **performance** depends crucially on “how well h reflects h^* ”!!
→ This is informally called the **informedness** or **quality** of h .
- For some search algorithms, like A^* , we can *prove* relationships between formal quality properties of h and search efficiency (mainly the number of expanded nodes).
- For other search algorithms, “it works well in practice” is often as good an analysis as one gets.

→ We will analyze in detail approximations to one particularly important heuristic function in planning: h^+ .

Heuristic Functions: Discussion, ctd.

“Search performance depends crucially on the informedness of h . . .”

Heuristic Functions: Discussion, ctd.

“Search performance depends crucially on the informedness of h . . .”

Any other property of h that search performance crucially depends on?

Heuristic Functions: Discussion, ctd.

“Search performance depends crucially on the informedness of h . . .”

Any other property of h that search performance crucially depends on?

“... and on the computational overhead of computing h !!”

Heuristic Functions: Discussion, ctd.

“Search performance depends crucially on the informedness of h . . .”

Any other property of h that search performance crucially depends on?

“... and on the computational overhead of computing h !!”

Extreme cases:

- $h = h^*$: Perfectly informed; computing it = solving the planning task in the first place.
- $h = 0$: No information at all; can be “computed” in constant time.

Heuristic Functions: Discussion, ctd.

“Search performance depends crucially on the informedness of h . . .”

Any other property of h that search performance crucially depends on?

“... and on the computational overhead of computing h !!”

Extreme cases:

- $h = h^*$: Perfectly informed; computing it = solving the planning task in the first place.
- $h = 0$: No information at all; can be “computed” in constant time.

→ Successful heuristic search requires a good trade-off between h ’s informedness and the computational overhead of computing it.

→ **This really is what research is all about!** Devise methods that yield good estimates at reasonable computational costs.

Properties of Heuristic Functions

Definition (Safe/Goal-Aware/Admissible/Consistent). Let Π be a planning task with state space $\Theta_{\Pi} = (S, L, c, T, I, S^G)$, and let h be a heuristic for Π . The heuristic is called:

- **safe** if $h^*(s) = \infty$ for all $s \in S$ with $h(s) = \infty$;
- **goal-aware** if $h(s) = 0$ for all goal states $s \in S^G$;
- **admissible** if $h(s) \leq h^*(s)$ for all $s \in s$;
- **consistent** if $h(s) \leq h(s') + c(a)$ for all transitions $s \xrightarrow{a} s'$.

Properties of Heuristic Functions

Definition (Safe/Goal-Aware/Admissible/Consistent). Let Π be a planning task with state space $\Theta_{\Pi} = (S, L, c, T, I, S^G)$, and let h be a heuristic for Π . The heuristic is called:

- **safe** if $h^*(s) = \infty$ for all $s \in S$ with $h(s) = \infty$;
- **goal-aware** if $h(s) = 0$ for all goal states $s \in S^G$;
- **admissible** if $h(s) \leq h^*(s)$ for all $s \in s$;
- **consistent** if $h(s) \leq h(s') + c(a)$ for all transitions $s \xrightarrow{a} s'$.

→ Relationships?

Properties of Heuristic Functions

Definition (Safe/Goal-Aware/Admissible/Consistent). Let Π be a planning task with state space $\Theta_{\Pi} = (S, L, c, T, I, S^G)$, and let h be a heuristic for Π . The heuristic is called:

- **safe** if $h^*(s) = \infty$ for all $s \in S$ with $h(s) = \infty$;
- **goal-aware** if $h(s) = 0$ for all goal states $s \in S^G$;
- **admissible** if $h(s) \leq h^*(s)$ for all $s \in S$;
- **consistent** if $h(s) \leq h(s') + c(a)$ for all transitions $s \xrightarrow{a} s'$.

→ Relationships?

Proposition. Let Π be a planning task with state space $\Theta_{\Pi} = (S, L, c, T, I, S^G)$, and let h be a heuristic for Π . If h is consistent and goal-aware, then h is admissible. If h is admissible, then h is goal-aware. If h is admissible, then h is safe. No other implications of this form hold.

Properties of Heuristic Functions

Definition (Safe/Goal-Aware/Admissible/Consistent). Let Π be a planning task with state space $\Theta_{\Pi} = (S, L, c, T, I, S^G)$, and let h be a heuristic for Π . The heuristic is called:

- **safe** if $h^*(s) = \infty$ for all $s \in S$ with $h(s) = \infty$;
- **goal-aware** if $h(s) = 0$ for all goal states $s \in S^G$;
- **admissible** if $h(s) \leq h^*(s)$ for all $s \in S$;
- **consistent** if $h(s) \leq h(s') + c(a)$ for all transitions $s \xrightarrow{a} s'$.

→ Relationships?

Proposition. Let Π be a planning task with state space $\Theta_{\Pi} = (S, L, c, T, I, S^G)$, and let h be a heuristic for Π . If h is consistent and goal-aware, then h is admissible. If h is admissible, then h is goal-aware. If h is admissible, then h is safe. No other implications of this form hold.

Proof. → **Exercise, perhaps.**

Agenda

- 1 Basics
- 2 Blind Systematic Search Algorithms
- 3 Heuristic Functions
- 4 Informed Systematic Search Algorithms**
- 5 Local Search Algorithms
- 6 Conclusion

Greedy Best-First Search

Greedy Best-First Search (with duplicate detection)

```

open := new priority queue ordered by ascending  $h(\text{state}(\sigma))$ 
open.insert(make-root-node(init()))
closed :=  $\emptyset$ 
while not open.empty():
     $\sigma := \text{open.pop-min()}$  /* get best state */
    if  $\text{state}(\sigma) \notin \text{closed}$ : /* check duplicates */
        closed := closed  $\cup \{\text{state}(\sigma)\}$  /* close state */
        if is-goal(state( $\sigma$ )): return extract-solution( $\sigma$ )
        for each  $(a, s') \in \text{succ}(\text{state}(\sigma))$ : /* expand state */
             $\sigma' := \text{make-node}(\sigma, a, s')$ 
            if  $h(\text{state}(\sigma')) < \infty$ : open.insert( $\sigma'$ )
return unsolvable

```

A*

A* (with duplicate detection and re-opening)

```


open := new priority queue ordered by ascending  $g(\text{state}(\sigma)) + h(\text{state}(\sigma))$ 
open.insert(make-root-node(init()))
closed :=  $\emptyset$ 
best-g :=  $\emptyset$  /* maps states to numbers */
while not open.empty():
     $\sigma := \text{open.pop-min}()$ 
    if  $\text{state}(\sigma) \notin \text{closed}$  or  $g(\sigma) < \text{best-g}(\text{state}(\sigma))$ :
        /* re-open if better g; note that all  $\sigma'$  with same state but worse g
           are behind  $\sigma$  in open, and will be skipped when their turn comes */
        closed := closed  $\cup$  {state( $\sigma$ )}
        best-g(state( $\sigma$ )) :=  $g(\sigma)$ 
        if is-goal(state( $\sigma$ )): return extract-solution( $\sigma$ )
        for each  $(a, s') \in \text{succ}(\text{state}(\sigma))$ :
             $\sigma' := \text{make-node}(\sigma, a, s')$ 
            if  $h(\text{state}(\sigma')) < \infty$ : open.insert( $\sigma'$ )
return unsolvable

```


A*: Example




A*: Example




A*: Example



A*: Example



A*: Example



A*: Example

