

Weighted A*

Weighted A* (with duplicate detection and re-opening)

```

open := new priority queue ordered by ascending  $g(\text{state}(\sigma)) + W * h(\text{state}(\sigma))$ 
open.insert(make-root-node(init()))
closed :=  $\emptyset$ 
best-g :=  $\emptyset$ 
while not open.empty():
     $\sigma := \text{open.pop-min}()$ 
    if  $\text{state}(\sigma) \notin \text{closed}$  or  $g(\sigma) < \text{best-g}(\text{state}(\sigma))$ :
        closed := closed  $\cup$  {state( $\sigma$ )}
        best-g(state( $\sigma$ )) :=  $g(\sigma)$ 
        if is-goal(state( $\sigma$ )): return extract-solution( $\sigma$ )
        for each  $(a, s') \in \text{succ}(\text{state}(\sigma))$ :
             $\sigma' := \text{make-node}(\sigma, a, s')$ 
            if  $h(\text{state}(\sigma')) < \infty$ : open.insert( $\sigma'$ )
return unsolvable

```

Weighted A*: Remarks

The **weight** $W \in \mathbb{R}_0^+$ is an **algorithm parameter**:

- For $W = 0$, weighted A* behaves like

Weighted A*: Remarks

The **weight** $W \in \mathbb{R}_0^+$ is an **algorithm parameter**:

- For $W = 0$, weighted A* behaves like uniform-cost search.
- For $W = 1$, weighted A* behaves like

Weighted A*: Remarks

The **weight** $W \in \mathbb{R}_0^+$ is an **algorithm parameter**:

- For $W = 0$, weighted A* behaves like uniform-cost search.
- For $W = 1$, weighted A* behaves like A*.
- For $W \rightarrow \infty$, weighted A* behaves like

Weighted A*: Remarks

The **weight** $W \in \mathbb{R}_0^+$ is an **algorithm parameter**:

- For $W = 0$, weighted A* behaves like uniform-cost search.
- For $W = 1$, weighted A* behaves like A*.
- For $W \rightarrow \infty$, weighted A* behaves like greedy best-first search.

Properties:

- For $W > 1$, weighted A* is **bounded suboptimal**: if h is admissible, then the solutions returned are at most a factor W more costly than the optimal ones.

Agenda

- 1 Basics
- 2 Blind Systematic Search Algorithms
- 3 Heuristic Functions
- 4 Informed Systematic Search Algorithms
- 5 Local Search Algorithms
- 6 Conclusion

Hill-Climbing

Hill-Climbing

```

 $\sigma := \text{make-root-node}(\text{init}())$ 
forever:
  if is-goal(state( $\sigma$ )):
    return extract-solution( $\sigma$ )
   $\Sigma' := \{ \text{make-node}(\sigma, a, s') \mid (a, s') \in \text{succ}(\text{state}(\sigma)) \}$ 
   $\sigma := \text{an element of } \Sigma' \text{ minimizing } h$  /* (random tie breaking) */

```

Hill-Climbing

Hill-Climbing

```

 $\sigma$  := make-root-node(init())
forever:
    if is-goal(state( $\sigma$ )):
        return extract-solution( $\sigma$ )
     $\Sigma'$  := { make-node( $\sigma, a, s'$ ) | ( $a, s'$ ) ∈ succ(state( $\sigma$ )) }
     $\sigma$  := an element of  $\Sigma'$  minimizing  $h$  /* (random tie breaking) */

```

Remarks:

- Makes sense only if $h(s) > 0$ for $s \notin S^G$.
- Is this complete or optimal?

Enforced Hill-Climbing

Enforced Hill-Climbing: Procedure *improve*

```

def improve( $\sigma_0$ ):
    queue := new fifo queue
    queue.push-back( $\sigma_0$ )
    closed :=  $\emptyset$ 
    while not queue.empty():
         $\sigma$  = queue.pop-front()
        if state( $\sigma$ )  $\notin$  closed:
            closed := closed  $\cup$  {state( $\sigma$ )}
            if  $h(\text{state}(\sigma)) < h(\text{state}(\sigma_0))$ : return  $\sigma$ 
            for each ( $a, s'$ )  $\in$  succ(state( $\sigma$ )):
                 $\sigma'$  := make-node( $\sigma, a, s'$ )
                queue.push-back( $\sigma'$ )

    fail
  
```

↪ Breadth-first search for state with strictly smaller h -value.

Enforced Hill-Climbing, ctd.

Enforced Hill-Climbing

```
 $\sigma$  := make-root-node(init())  
while not is-goal(state( $\sigma$ )):  
     $\sigma$  := improve( $\sigma$ )  
return extract-solution( $\sigma$ )
```

Enforced Hill-Climbing, ctd.

Enforced Hill-Climbing

```
 $\sigma := \text{make-root-node}(\text{init}())$   
while not  $\text{is-goal}(\text{state}(\sigma))$ :  
     $\sigma := \text{improve}(\sigma)$   
return  $\text{extract-solution}(\sigma)$ 
```

Remarks:

- Makes sense only if $h(s) > 0$ for $s \notin S^G$.
- Is this optimal?

Agenda

- 1 Basics
- 2 Blind Systematic Search Algorithms
- 3 Heuristic Functions
- 4 Informed Systematic Search Algorithms
- 5 Local Search Algorithms
- 6 Conclusion

AI Planning for Autonomy

3. Introduction to Planning

How to Describe Arbitrary Search Problems

Chris Ewin & Tim Miller




THE UNIVERSITY OF
MELBOURNE

With slides by Nir Lipovetsky

Beating Kasparov is great . . .



Beating Kasparov is great . . . but how to play Mario?



- You (and your brother/sister/little nephew) are better than Deep Blue at **everything** - except playing Chess.
- Is that (artificial) 'Intelligence'?

→ How to build machines that automatically solve **new** problems?

Planning: Motivation

How to develop systems or 'agents'
that can make decisions on their own?

Autonomous Behavior in AI

The key problem is to select **the action to do next**. This is the so-called **control problem**. Three approaches to this problem:

- **Programming-based:** Specify control by hand
- **Learning-based:** Learn control from experience
- **Model-based:** Specify problem by hand, derive control automatically

→ Approaches not orthogonal; successes and limitations in each ...

→ Different **models** yield different types of **controllers** ...

Programming-Based Approach

→ Control specified by programmer, e.g.:

- If Mario finds no danger, then run...
- If danger appears and Mario is big, jump and kill ...
- ...

- **Advantage**: domain-knowledge easy to express
- **Disadvantage**: cannot deal with situations not anticipated by programmer

Learning-Based Approach

→ Learns a controller from experience or through simulation:

■ **Unsupervised** (Reinforcement Learning):

- penalize Mario each time that 'dies'
- reward agent each time opponent 'dies' and level is finished, . . .

■ **Supervised** (Classification)

- learn to classify actions into good or bad from info provided by teacher

■ **Evolutionary**:

- from pool of possible controllers: try them out, select the ones that do best, and mutate and recombine for a number of iterations, keeping best

■ **Advantage**: does not require much knowledge in principle

■ **Disadvantage**: in practice, hard to know which features to learn, and is slow

General Problem Solving

Ambition: Write **one** program that can solve **all** problems.

→ Write $X \in \{\text{algorithms}\} : \text{for all } Y \in \{\text{'problems'}\} : X \text{'solves'} Y$

→ What is a 'problem'? What does it mean to 'solve' it?

Ambition 2.0: Write one program that can solve **a large class of problems**

Ambition 3.0: Write one program that can solve a large class of problems **effectively**


(some new problem) \leadsto (describe problem → use off-the-shelf solver) \leadsto (solution competitive with a human-made specialized program)

→ **Beat humans at coming up with clever solution methods!**

(Link: GPS started on 1959)

Model-Based Approach / General Problem Solving

- specify model for problem: **actions, initial situation, goals, and sensors**
- let a solver compute controller automatically



Model-Based Approach / General Problem Solving

→ Advantage:

- **Powerful**: In some applications generality is absolutely necessary
- **Quick**: Rapid prototyping. 10s lines of problem description vs. 1000s lines of C++ code. (Language generation!)
- **Flexible & Clear**: Adapt/maintain the description.
- **Intelligent & domain-independent**: Determines automatically how to solve a complex problem effectively! (The ultimate goal, no?!)

→ Disadvantage: need a model; computationally intractable

- **Efficiency loss**: Without any domain-specific knowledge about Chess, you don't beat Kasparov . . .

→ Trade-off between 'automatic and general' vs. 'manualwork but effective'

Model-based approach to intelligent behavior called **Planning** in AI

How to make fully automatic algorithms effective?

Agenda

1 Models

2 Languages

3 Complexity

4 Computational Approaches

5 IPC

6 Conclusion

Basic State Model: Classical Planning

Ambition:

Write one program that can solve all classical search problems.

State Model:

- finite and discrete state space S
- a **known initial state** $s_0 \in S$
- a set $S_G \subseteq S$ of goal states
- actions $A(s) \subseteq A$ applicable in each $s \in S$
- a **deterministic transition function** $s' = f(a, s)$ for $a \in A(s)$
- positive **action costs** $c(a, s)$

→ A **solution** is a sequence of applicable actions that maps s_0 into S_G , and it is **optimal** if it minimizes **sum of action costs** (e.g., # of steps)

→ Different **models** and **controllers** obtained by relaxing assumptions in **blue** ...

Uncertainty but No Feedback: Conformant Planning

- finite and discrete state space S
- a **set of possible initial state** $S_0 \in S$
- a set $S_G \subseteq S$ of goal states
- actions $A(s) \subseteq A$ applicable in each $s \in S$
- a **non-deterministic** transition function $F(a, s) \subseteq S$ for $a \in A(s)$
- uniform action costs $c(a, s)$

→ A **solution** is still an **action sequence** but must achieve the goal for **any possible initial state and transition**

→ More complex than **classical planning**, verifying that a plan is **conformant** intractable in the worst case; but special case of **planning with partial observability**

Planning with Markov Decision Processes

MDPs are **fully observable, probabilistic** state models:

- a state space S
- initial state $s_0 \in S$
- a set $G \subseteq S$ of goal states
- actions $A(s) \subseteq A$ applicable in each state $s \in S$
- **transition probabilities** $P_a(s'|s)$ for $s \in S$ and $a \in A(s)$
- action costs $c(a, s) > 0$

→ **Solutions** are **functions (policies)** mapping states into actions

→ **Optimal** solutions minimize **expected cost** to goal

Partially Observable MDPs (POMDPs)

POMDPs are **partially observable, probabilistic** state models:

- states $s \in S$
- actions $A(s) \subseteq A$
- transition probabilities $P_a(s'|s)$ for $s \in S$ and $a \in A(s)$
- initial **belief state** b_0
- final **belief state** b_f
- **sensor model** given by probabilities $P_a(o|s)$, $o \in Obs$


→ **Belief states** are probability distributions over S

→ **Solutions** are policies that map belief states into actions

→ **Optimal** policies minimize **expected** cost to go from b_0 to G

Example

Agent **A** must reach **G**, moving one cell at a time in **known** map



- If actions deterministic and initial location known, planning problem is **classical**
- If actions stochastic and location observable, problem is an **MDP**
- If actions stochastic and location partially observable, problem is a **POMDP**

Different combinations of uncertainty and feedback: three problems, three models

Agenda

1 Models

2 Languages

3 Complexity

4 Computational Approaches

5 IPC

6 Conclusion

Models, Languages, and Solvers

- A **planner** is a **solver over a class of models**; it takes a model description, and computes the corresponding controller

$$Model \implies \boxed{Planner} \implies Controller$$

- Many models, many solution forms: uncertainty, feedback, costs, ...
- Models described in suitable **planning languages** (Strips, PDDL, PPDDL, ...) where **states** represent interpretations over the language.

A Basic Language for Classical Planning: Strips

- A **problem** in STRIPS is a tuple $P = \langle F, O, I, G \rangle$:
 - F stands for set of all **atoms** (boolean vars)
 - O stands for set of all **operators** (actions)
 - $I \subseteq F$ stands for **initial situation**
 - $G \subseteq F$ stands for **goal situation**
- Operators $o \in O$ **represented** by
 - the **Add** list $Add(o) \subseteq F$
 - the **Delete** list $Del(o) \subseteq F$
 - the **Precondition** list $Pre(o) \subseteq F$

From Language to Models (STRIPS Semantics)


A STRIPS problem $P = \langle F, O, I, G \rangle$ determines **state model** $S(P)$ where

- the states $s \in S$ are **collections of atoms** from F . $S = 2^F$
- the initial state s_0 is I
- the goal states s are such that $G \subseteq s$
- the actions a in $A(s)$ are ops in O s.t. $Prec(a) \subseteq s$
- the next state is $s' = s - Del(a) + Add(a)$
- action costs $c(a, s)$ are all 1

→ (Optimal) **Solution** of P is (optimal) **solution** of $S(P)$

→ Slight language extensions often convenient: **negation, conditional effects, non-boolean variables**; some required for describing richer models (costs, probabilities, ...).

(Oh no it's) The Blockworld



- **Propositions:** $on(x, y)$, $onTable(x)$, $clear(x)$, $holding(x)$, $armEmpty()$.
- **Initial state:** $\{onTable(E), clear(E), \dots, onTable(C), on(D, C), clear(D), armEmpty()\}$.
- **Goal:** $\{on(E, C), on(C, A), on(B, D)\}$.
- **Actions:** $stack(x, y)$, $unstack(x, y)$, $putdown(x)$, $pickup(x)$.
- $stack(x, y)?$